


Chapter 1: Introduction

Reading: Chapter 1.1-1.6





Today's Lecture

- Reading 1.1-1.3
 - What Operating Systems do
 - Computer System Organization





Operating System: What comes to your mind?



chromebook





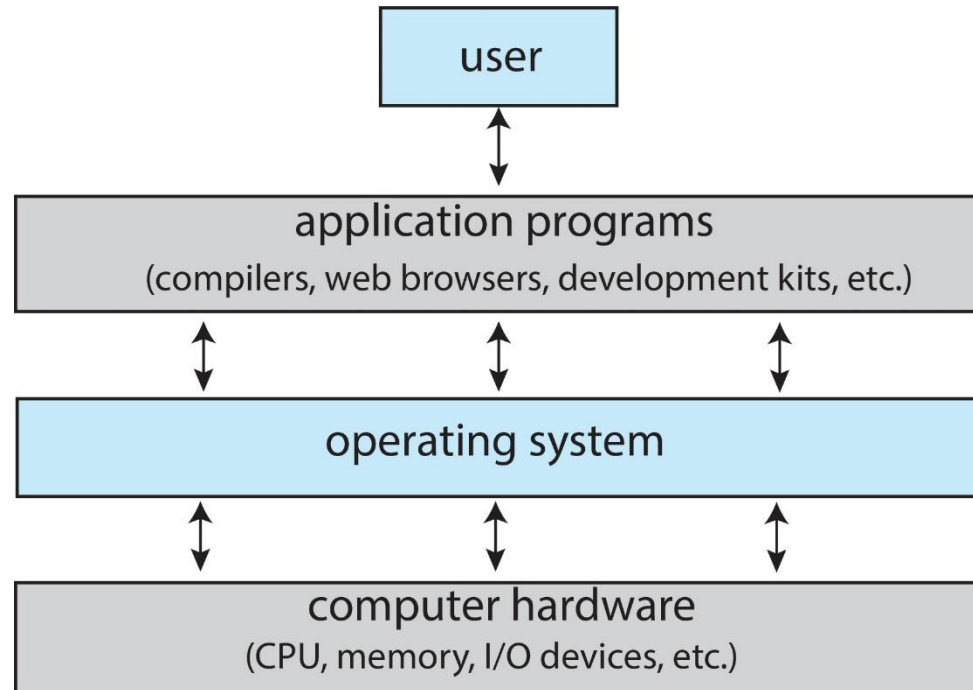
Operating System

- ❑ Software that manages computer hardware
- ❑ Acts as an intermediary between user and hardware
- ❑ Provides an environment for software applications





Components of a Computer System





What Operating Systems Do

- OS is a **resource allocator**:
 - Allocates resources (e.g., CPU time, memory space, disk space, I/O devices) to specific programs *efficiently* and *fairly*

- OS is a **control program**:
 - Controls the execution of user programs to prevent errors and improper use of the computer
 - Controls I/O devices






Operating System Definition

- An operating system consists of
 - **kernel** – core of OS, running at all times on the computer
 - **system programs** (ships with the OS, but not part of the kernel)
 - ▶ file management programs, device drivers, GUI systems, text editors, compilers, debuggers, etc.
- **Application programs** are all programs not associated with the OS
 - Web browsers, word processors, database systems, video games, etc.





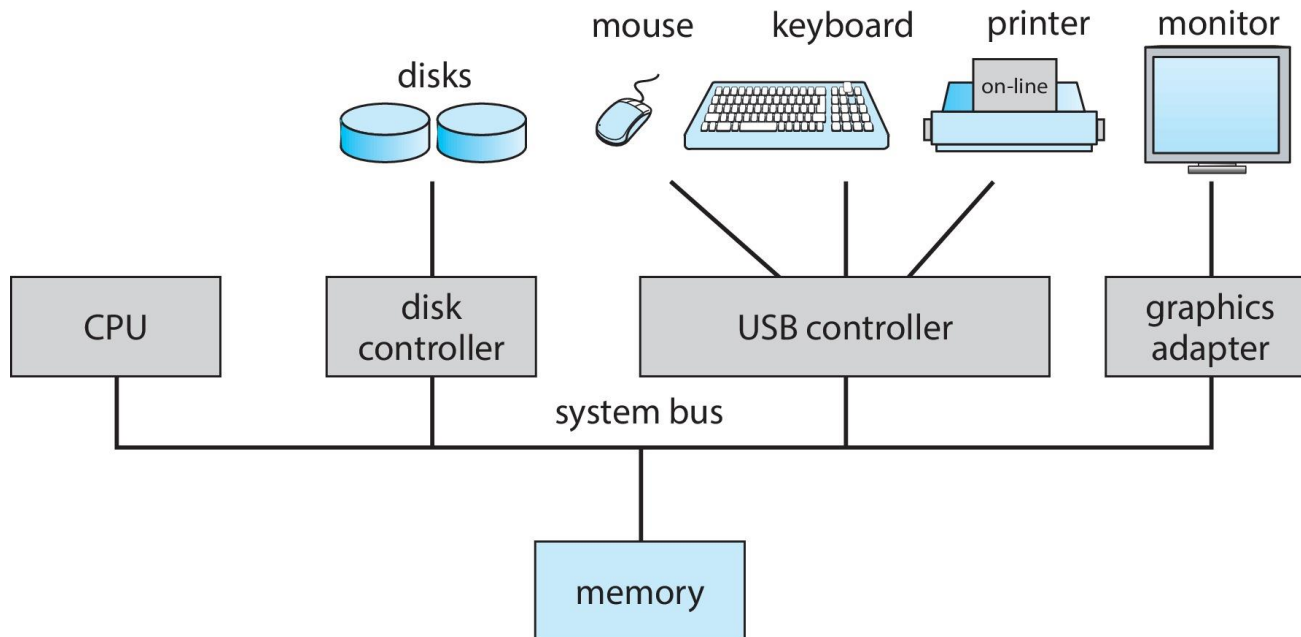
PC Motherboard





Computer System Organization

- One or more CPUs and device controllers connect through common **bus** providing access to shared memory
 - Each device controller is in charge of a particular device type
- CPUs and device controllers can execute in parallel, competing for memory cycles





I/O Operations

- ❑ OS has a **device driver** for each device controller
- ❑ Each device controller has a local buffer
 - ❑ Device controller moves data between its local buffer and the devices it controls
- ❑ Device driver initiates I/O on device controller
 - ❑ In case of output, CPU moves output data from main memory to local buffer of device controller
- ❑ Device controller informs CPU that it has finished its I/O operation by causing an **interrupt**
 - ❑ In case of input, CPU moves input data from local buffer of device controller to main memory





Interrupts

- ❑ Operating systems are **interrupt driven**
- ❑ An interrupt may be triggered by hardware (e.g., I/O device) or software
 - ❑ A software-generated interrupt is called a **trap** (or **exception**)
 - ❑ A **trap** is caused either by an error or a request from a user program
- ❑ Interrupt transfers control to the appropriate interrupt service routine (ISR) generally through the **interrupt vector**, which contains the addresses of all the ISRs
- ❑ Upon an interrupt, OS preserves the state of the CPU by storing registers and the program counter so that the CPU can resume the interrupted computation after servicing the interrupt
- ❑ Review Animation 1.2: Interrupt driver I/O cycle





Interrupt Vector Table


vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts

Figure 1.5 Intel processor event-vector table.



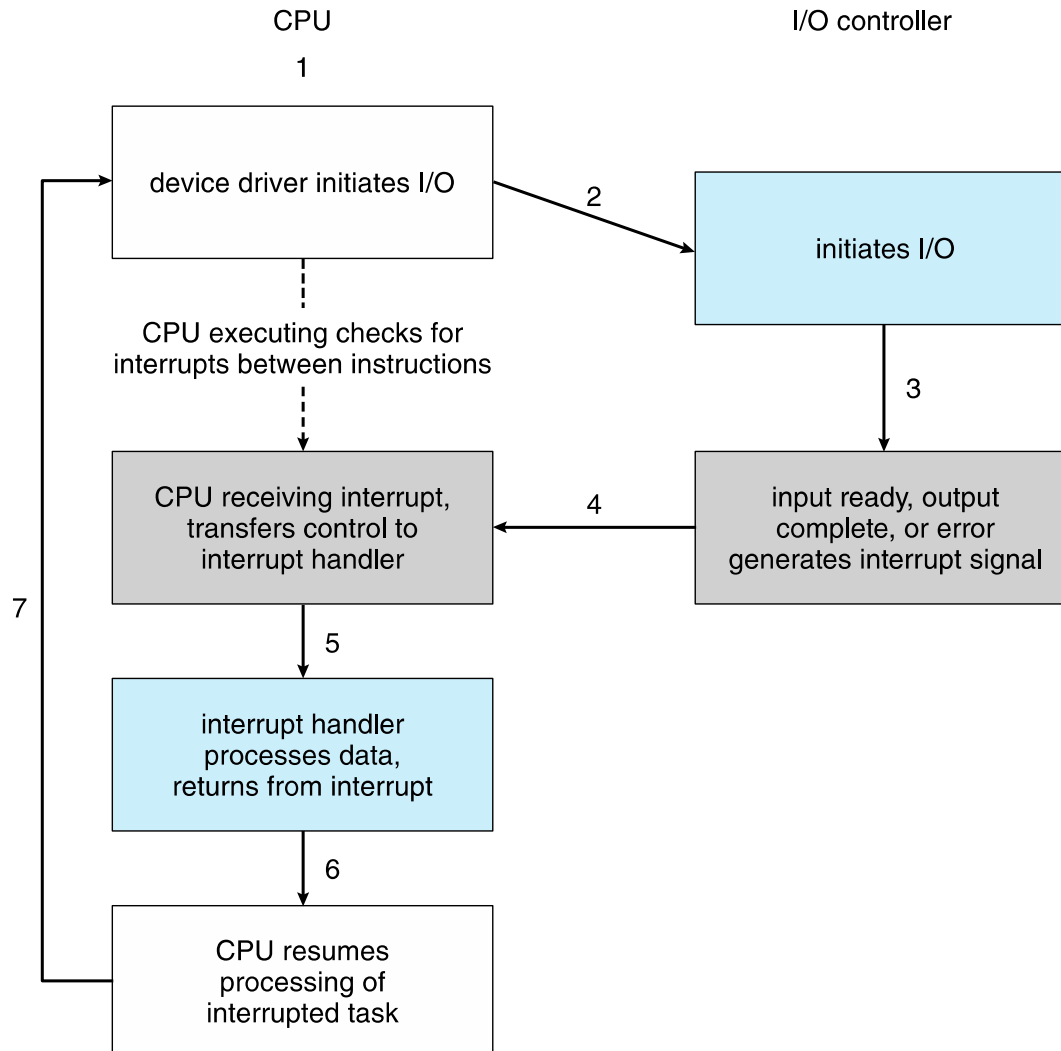


Interrupt Timeline





Interrupt-driven I/O Cycle





Multiprocessor system




Figure 1.8 Symmetric multiprocessing architecture.





Multicore system





Figure 1.9 A dual-core design with two cores on the same chip.





Storage-Device Hierarchy





Storage Structure

- Main memory – only large storage media that the CPU can access directly
 - Typically **volatile**, i.e., loses contents when power is removed
 - Any programs must be first loaded into main memory to run
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
 - **Mechanical: Hard Disk Drives (HDD)**
 - **Electrical: nonvolatile memory (NVM)** devices– faster than hard disks
 - ▶ Various technologies: flash memory, solid-state drive (SSD)





A short quiz

- ❑ Interrupts are triggered only by hardware
 - ❑ True
 - ❑ False





OS Operations

- OS is **interrupt driven**
 - Hardware interrupt generated by one of the devices
 - Software interrupt (**trap** or **exception**):
 - ▶ Error (division by zero, invalid memory access)
 - ▶ Request for operating system service – **system call**





Today's Lecture

- Reading 1.4-1.6
 - Operating-System Operations
 - Resource Management
 - Security and Protection





I/O Structure

- ❑ Drawbacks of interrupt-driven I/O
 - ❑ Transfer data one byte at a time
 - ❑ CPU involved in data transfer between memory and device controller
 - ❑ Incurs high overhead when used for bulk data movement (e.g., disk I/O)
- ❑ **Direct Memory Access (DMA)** – DMA controller tells the device controller to transfer a block of data directly between its buffer storage and memory without CPU intervention
 - ❑ CPU can do other work while data transfer is in progress
 - ❑ DMA controller interrupts the CPU when entire block is transferred





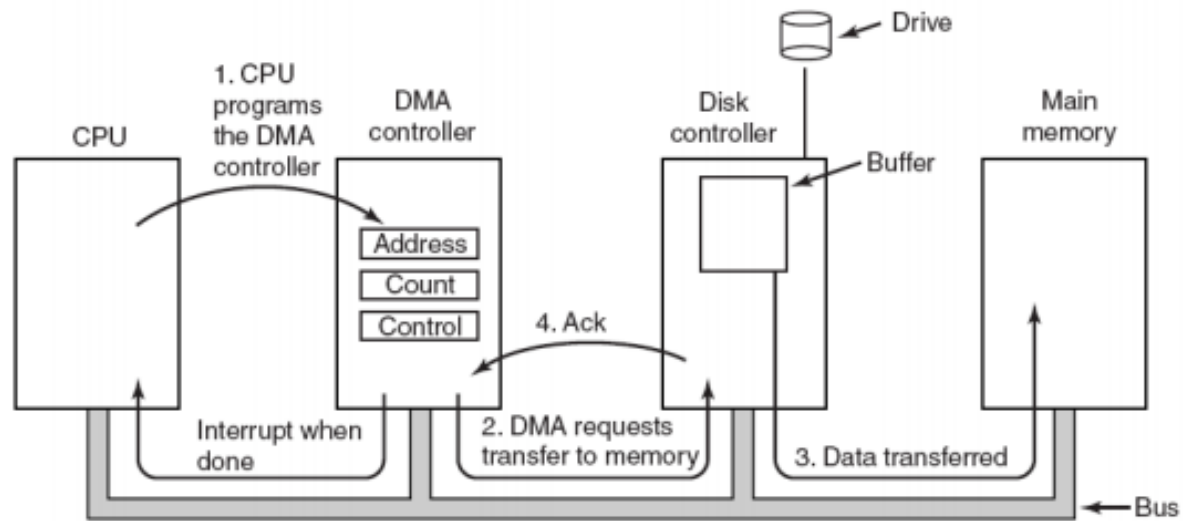
How a Modern Computer Works

- Review Animation 1.3





DMA




Operation of a DMA transfer





Multiprogramming

- ❑ **Multiprogramming** needed for efficiency
 - ❑ Single program cannot keep CPU and I/O devices busy at all times
 - ❑ OS keeps several **processes** in memory simultaneously so CPU always has a process to execute
 - ▶ A program in execution is called a **process**
 - ❑ One process is selected and run via **CPU scheduling**
 - ❑ When a process has to wait (e.g., for I/O completion), OS switches to another process



Memory layout





Exercise

Consider two programs:

- Program A: computation (100ms), then I/O on device 1 (200ms), then computation (50ms)
- Program B: computation (50ms), then I/O on device 2 (200ms), then computation(100ms)



How long it will take to complete the two programs when

- (1) Monoprogramming (i.e., only one program is kept in memory) is used?
- (2) Multiprogramming is used?





Monoprogramming




Total time = 700ms





Multiprogramming



Total time = 450ms






Your turn (10 mins)

Consider three programs:

- Program A: computation (50ms), then I/O on device 1 (200ms), then computation (50ms)
- Program B: computation (50ms), then I/O on device 2 (100ms), then computation(10ms)
- Program C: computation (50ms), then I/O on device 2 (100ms)

- Monoprogramming: Total time: 610ms
- Multiprogramming: Total time: 300ms







Multitasking

- A multiprogramming OS focuses on efficient use of CPU and I/O devices
- **Multitasking** (**timesharing**) is a logical extension of multiprogramming in which users can interact with each process while it is running, creating **interactive** computing
 - OS rapidly switches between users' processes to provide fast response time to users
 - Each user is given the impression that the entire computer system is dedicated to his/her use





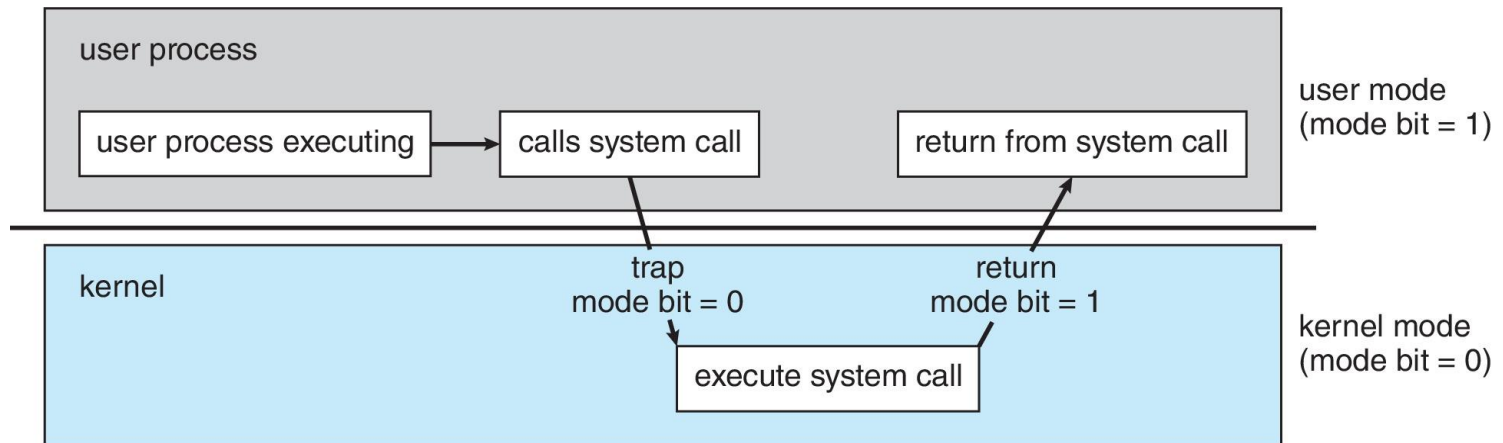
Dual-Mode Operation

- **Dual-mode** operation allows OS to protect itself and user programs
 - Ensures that an incorrect (or malicious) program cannot cause other programs or the OS to execute incorrectly
- CPU can operate in two modes: **User mode** (1) and **kernel mode** (0)
 - **Mode bit** provided by hardware, distinguishes when system is running user code or kernel code
- Some instructions are designated as **privileged instructions**, only executable in kernel mode
 - Such instructions may cause harm to system resources (I/O devices, memory, files) and user/OS code
 - E.g., clear memory, access I/O device
- User programs execute in user mode. When a trap or interrupt occurs, OS puts CPU in kernel mode
- OS puts CPU in user mode before passing control to a user program





Transition from User to Kernel Mode



System call changes mode to kernel mode, return from call resets mode to user mode





Timer

- ❑ To prevent a process from hogging resources, a timer is set to interrupt the CPU after some time period
- ❑ A timer is implemented by the physical clock and a counter
 - ❑ OS sets the counter (privileged instruction)
 - ❑ Counter is decremented every time the clock ticks
 - ❑ When counter reaches zero, generate an interrupt
 - ❑ OS sets up timer before scheduling a process in order to regain control





Process Management

- A **process** is a program in execution; it is a unit of work in an OS
- Process needs resources to accomplish its task
 - CPU time, memory, I/O devices, files
 - Process termination requires OS to reclaim any reusable resources
- OS activities in connection with process management:
 - Creating and deleting both user and system processes
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization
 - Providing mechanisms for process communication
 - Providing mechanisms for deadlock handling





Memory Management

- To execute a program:
 - All or part of the instructions must be in memory
 - All or part of the data needed by the program must be in memory
- OS keeps multiple programs in memory to improve CPU utilization and speed of computer response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed





File-System Management

- OS provides a uniform, logical view of information storage
 - OS defines a logical storage unit – **file**
 - OS maps files onto physical media

- File-system management activities
 - Creating and deleting files and directories
 - Supporting primitives for manipulating files and directories
 - Mapping files onto secondary storage
 - Access control to determine who can access what





Secondary-Storage Management

- Secondary storage (e.g., hard disk drives, nonvolatile memory devices) used to store programs and data that do not fit in main memory
- OS activities
 - Mounting and unmounting file systems
 - Free-space management
 - Storage allocation
 - Disk scheduling
 - Disk partitioning





Cache Management

- **Caching** – information in use is copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache is smaller than storage being cached
 - Cache management an important design problem
 - Cache size and replacement policy





Characteristics of Various Types of Storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid-state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25-0.5	0.5-25	80-250	25,000-50,000	5,000,000
Bandwidth (MB/sec)	20,000-100,000	5,000-10,000	1,000-5,000	500	20-150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape





Protection & Security, Networking

- **Protection** – mechanism for controlling access of processes or users to resources of computer system
 - OS distinguishes between authorized and unauthorized usage of resources
- **Security** – defense of the system against external attacks
 - E.g., denial-of-service attacks, worms, viruses, identity theft, theft of service
- **Networking** – OS support network protocols (e.g., TCP/IP)

