


Chapter 2: Operating-System Structures

Reading: Chapter 2.1-2.3





Operating System Services

- ❑ An operating system provides services to programs and users
- ❑ One set of OS services provides functions helpful to the user
 - ❑ **User interface** – **command-line interface (or shell)**, **graphical user interface (GUI)**, **touch-screen**
 - ❑ **Program execution** - OS must be able to load a program into memory, run that program, and end program execution
 - ❑ **I/O operations** - A running program may require I/O, which may involve a file or an I/O device
 - ❑ **File-system manipulation** (e.g., read/write files and directories, create/delete files and directories)
 - ❑ **Communications** – Processes may exchange information via shared memory or through message passing
 - ❑ **Error detection**
 - ▶ Errors may occur in hardware or in user program
 - ▶ For each type of error, OS should take the appropriate action to ensure correct and consistent computing






Operating System Services (Cont.)

- Another set of OS functions ensure the efficient operation of the system itself
 - **Resource allocation** - When multiple processes run concurrently, resources must be allocated to each of them
 - **Logging** - To keep track of which users use how much and what kinds of computer resources
 - **Protection and security**
 - ▶ **Protection** involves ensuring that all access to system resources is controlled
 - ▶ **Security** is to defend a system from external attacks





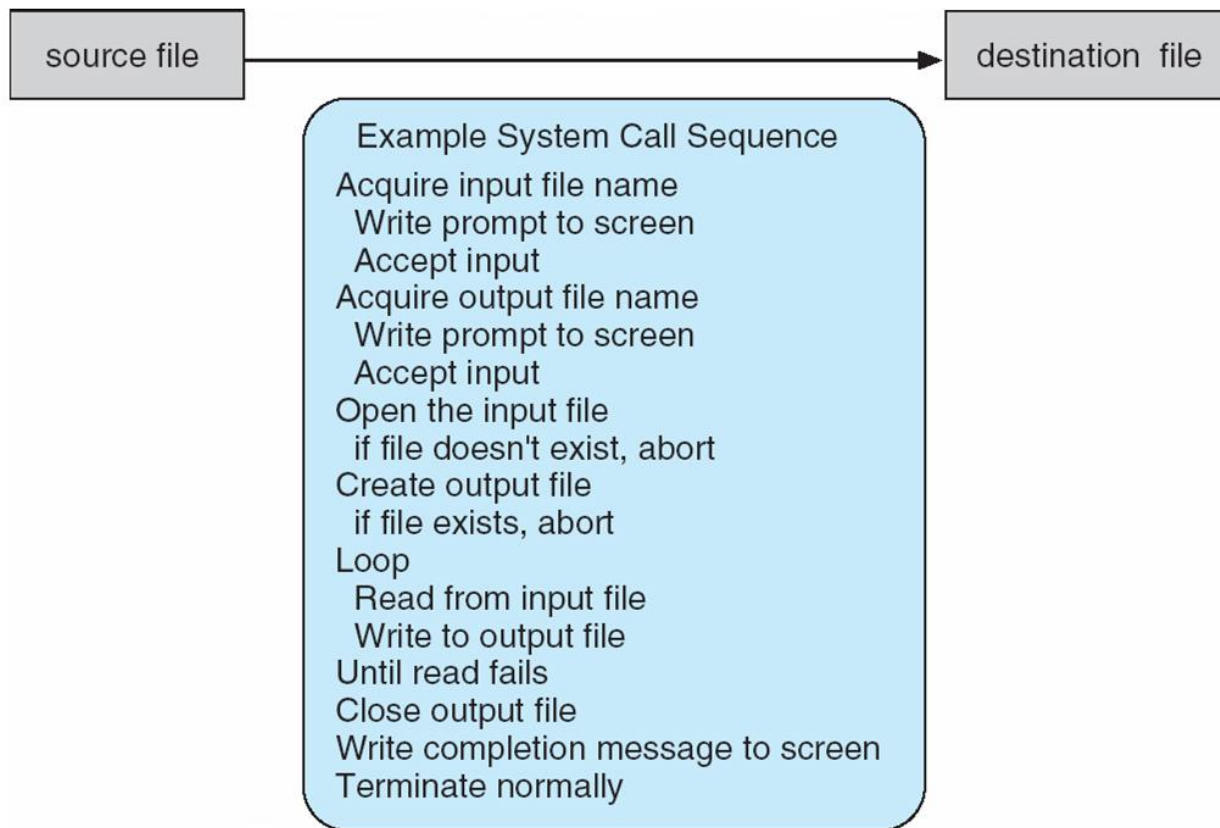
A View of Operating System Services





System Calls

- System calls provide an interface to the services provided by the OS
 - Generally available as routines written in C or C++
- Example of system calls



System call sequence to copy the contents of one file to another file





Example of Standard API

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the man page by invoking the command

`man read`

on the command line. A description of this API appears below:

```
#include <unistd.h>

ssize_t  read(int fd, void *buf, size_t count)
```

return value	function name	parameters
--------------	---------------	------------

→ Type “man 2 read” to look up read in section 2 of man page

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.





Application Programming Interface

- Programmers mostly access systems calls via a high-level **Application Programming Interface (API)** rather than direct system call use
 - The API specifies a set of library functions available to application programmers
 - The library functions typically invoke system calls
- Three most common APIs
 - Windows API for Windows systems
 - POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and macOS)
 - Java API for programs that run on the Java virtual machine (JVM)

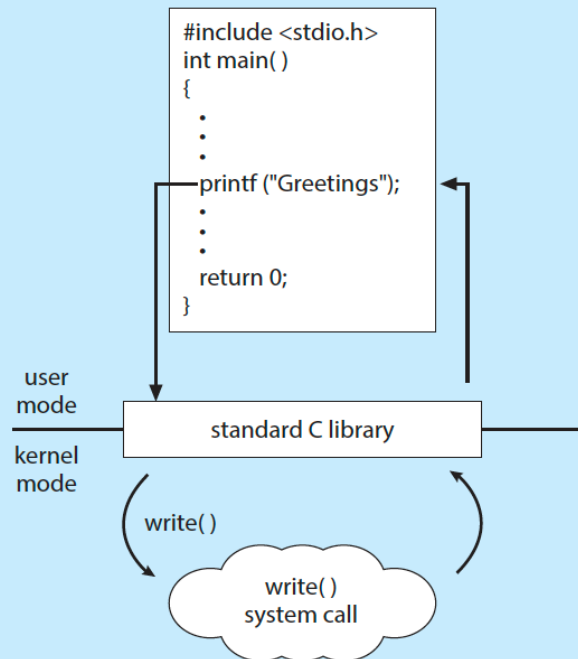




Standard C Library Example

THE STANDARD C LIBRARY

The standard C library provides a portion of the system-call interface for many versions of UNIX and Linux. As an example, let's assume a C program invokes the `printf()` statement. The C library intercepts this call and invokes the necessary system call (or calls) in the operating system—in this instance, the `write()` system call. The C library takes the value returned by `write()` and passes it back to the user program:



A C program invoking `printf()` library call, which invokes `write()` system call






System Call Implementation

- When a user program makes a system call, it triggers a trap
- Typically, a number is associated with each system call
 - The **run-time environment** of a given programming language provides a **system-call interface** that maintains a table indexed according to these numbers
 - Table stores addresses of system call implementations
- The system-call interface invokes the intended system call in OS kernel and returns status of the system call and any return values
- The caller knows nothing about how the system call is implemented
 - Just needs to obey API and understand what OS will do as a result of the call





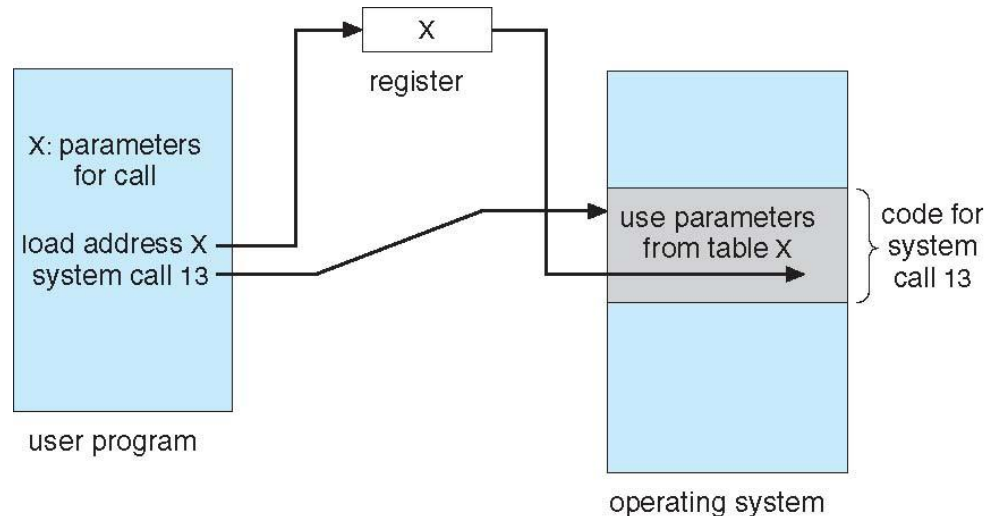
Example: User Application Invoking `open()` System Call





System Call Parameter Passing

- Three general methods are used to pass parameters to the OS
 - Method 1: Pass the parameters in registers
 - ▶ There may be more parameters than registers
 - Method 2: Parameters stored in a block in memory, and address of block passed as a parameter in a register



- Method 3: Parameters **pushed** onto the **stack** by the program and **popped** off the stack by the OS
- Methods 2 and 3 do not limit the number or length of parameters being passed

